


Beyond Vibe Coding

Architecting Agentic Workflows for Production Engineering



The Default Path is the Shortest Path (to Technical Debt)

- AI agents act like talented but **reckless** junior engineers.
- They **optimize** for **speed** over process.
- They produce code, **declare victory**, and move on.



> "Task complete. I have implemented the feature.
I will write the tests later."



Bolting the Senior Engineer Back On

- Senior engineering is mostly the parts that don't show up in the diff: surfacing assumptions, scoping, leaving evidence of correctness, and verifying.

~~Specs~~

~~Tests~~

~~Security Review~~

The Solution: Agentic Workflows (Process over Prose).

Instead of writing a 2,000-word prompt about best practices, we encode hard exit criteria into the agent's environment.

Specs: Surfacing Assumptions

Tests: Evidence of Correctness

Security Review: Verification

✓ Verifiable PR: Done

Four Philosophies of AI Scaffolding



VS Code Plan

The Native Draftsman
(Built-in planning)

Matt Pocock's Skills

The Interrogator
(Composable toolkit)



Obra's Superpowers

The Autonomous
Assembly Line
(Sub-agent pipeline)

Addy Osmani's Agent Skills

The Governance Board
(Opinionated full SDLC)




The choice is not about syntax; it is about matching an architectural paradigm to your team's culture.

The Native Draftsman: VS Code Plan

Core Philosophy: Native integration and built-in task separation.

- **Deep Research:** The agent deploys sub-agents to read the codebase and external docs before writing code.
- **Artifact Generation:** It generates a highly structured plan.md detailing todo lists and affected files.
- **Custom Personas:** Engineers can supply custom instructions to create planning personas tailored to their specific tech stack.



The VS Code Handoff Loop


By separating the 'Plan' harness from the execution 'Agent,' VS Code prevents the AI from rushing to modify files before the architecture is settled.



The Interrogator: Matt Pocock's Skills

Core Philosophy: Fix requirements first, keep tools composable, and respect the 'Smart Zone' limit.

- ✂ **The Swiss Army Knife:** Designed to be small, adaptable modules rather than an end-to-end dictated process.
- ✂ **Stateful Domain Building:** Updates CONTEXT.md and ADRs inline during exploration.
- ✂ **The 140k Token Rule:** Breaks work down so each execution ticket fits neatly into the LLM's reliable 'smart zone'.



The grill-with-docs Funnel


The agent refuses to proceed until every **branch** of the design tree is resolved and a **shared understanding is confirmed**.



The Autonomous Assembly Line: Obra's Superpowers


Core Philosophy: Deep upfront reasoning followed by a strict, hands-off execution pipeline.

- ✂ **Micro-Tasking:** Work is broken into 2-5 minute bite-sized tasks.
- ✂ **Strict RED-GREEN-REFACTOR:** Enforces TDD. Any implementation code written before a failing test is immediately deleted.
- ✂ **Two-Stage Review:** Code is verified first for spec compliance, then for quality constraints (DRY/YAGNI).



Sub-Agent Driven Development


Parallel tasks are given fresh, clean context windows and isolated workspaces, preventing agents from clobbering each other's code.



The Governance Board: Addy Osmani's Agent Skills


Core Philosophy: Encode the entire Software Development Lifecycle (SDLC) with adversarial checkpoints.

- **Process over Prose:** 24 workflows encoding Google-scale engineering norms (Hyrum's Law, Chesterton's Fence).
- **Verification is Non-Negotiable:** Tests, builds, and runtime checks (like Chrome DevTools profiling) act as hard exit criteria.
- **Parallel Personas:** Deploying specialized sub-agents (security auditor, performance tester, QA) simultaneously before shipping.




The Anti-Rationalization Wall

Pre-written rebuttals to lies the agent hasn't yet told. It prevents the AI from justifying dangerous shortcuts.



The Architectural Landscape

There is no 'best' framework—only the right balance of control and autonomy for your team's culture.



The Workflow Diagnostic Matrix

	VS Code Plan	Matt Pocock	Obra Superpowers	Agent Skills
Core Philosophy	Built-in Draftsman	Interrogator	Assembly Line	Governance Board
Signature Mechanism	`plan.md` Handoff	`grill-me` Funnel	Git Worktrees	Anti-Rationalization
Autonomy Level	Guided	High Human Input	Max Autonomy	Gated / Checked
Best For...	Native editor integrations	Unclear requirements	Large, exploratory tasks	Standardizing team SDLC

The Architectural Landscape

Choose a Philosophy, Not a Prompt.

The senior-engineer parts of the job—surfacing assumptions, sizing changes, and leaving evidence—are no longer optional.

Whether you compose a toolkit or deploy a pipeline, the goal remains the same:
Process over prose. Verification over vibes.

